

CSE 517: Homework #4

Structured Learning with Feature-Rich Models

Due Wed Mar 1 2017 11:59pm

Hi Class! So, the remaining two assignments will be all about structured learning and inference. In Part I (HW #4), you will think about how to construct perceptron algorithms for structured learning. (No worries if you haven't formally learned the perceptron algorithm before. This assignment is self-contained!) To convert a flat perceptron algorithm into a structured perceptron one, you will need to apply the same intuition we used for converting MaxEnt models into CRF models. In the process, you will also think more carefully about the internals of feature-rich models. In Part II (HW #5), you will play with encoder-decoder RNN architecture (LSTMs and GRUs!) that has brought a lot of interesting results in recent literature. Hope this will be fun!

Submission instructions As usual, submit 2 files on Canvas — **Code** (HW5.tgz) and **Report** (HW5.pdf). The usual submission guideline applies, which we omit now for brevity.

1 Named-Entity Recognition (NER) with Structured Perceptrons

Named entities are phrases that contain the names of persons, organization, locations, etc [3]. For example, in the following NER-tagged sentence[1],

[ORG U.N.] official [PER Ekeus] heads for [LOC Baghdad] .

U.N is an *organization*, Ekeus is a *person*, and Baghdad is a *location*.

NER as sequence tagging: NER can be formulated as a sequence tagging problem using **BIO encoding**, where the **B**eginning word of a named entity is marked with '**B**', the following words of a named entity is marked with '**I**' (i.e., inside) and words that are not part of an entity is marked with '**O**' (i.e., outside). These labels are further marked with the types of named entities (e.g., ORG, PER, LOC). See section 1.2 for a concrete example of the tagging scheme.

Feature Vectors: You are free to design the feature vector however you like. You can start with feature templates similar to those presented in page 14 of LogLinear.pdf slides, or you can look up Table 1 – 3 of [4]. See J&M textbook Ch#22.1 if you like to see yet additional examples. But it's also fine if you'd rather be creative and come up with your own features without reading all these references.

1.1 Sequential Tagging with Structured Perceptrons

The pseudo code in **Algorithm 1** on the next page provides the recipe for the structured perceptron learning [2], which is almost identical to perceptron algorithm for simple classification, except that \mathbf{y} here is a sequence of random variables ($\mathbf{y} = y_1, \dots, y_n$) as opposed to one random variable.

The Equation (1) and (2) summarize the perceptron algorithm — you predict $\tilde{\mathbf{y}}$ of a training example \mathbf{x}^i using the current weight vector \mathbf{w} (Equation (1)). If the current weight vector leads to an incorrect prediction,

then update the parameters with the difference between the feature vector of the correct prediction and the feature vector of the incorrect prediction (Equation (2)). While simple, perceptron algorithm often leads to performance that is almost as good as state-of-the-art if given a good feature vector.

The reason why we take the averaged weight vector $\bar{\mathbf{w}}$ is because it gives a bit of a regularization effect (i.e., counteracting overfitting). But sometimes people take the last vector \mathbf{w} directly as the final learned parameters.

Remember that you will need to factorize the global feature vector $\Phi(\mathbf{x}^i, \mathbf{y})$ as a summed vector over local feature vectors, for example,

$$\Phi(\mathbf{x}^i, \mathbf{y}) = \sum_k \phi(\mathbf{x}^i, k, \mathbf{y}_k, \mathbf{y}_{k-1})$$

or even

$$\Phi(\mathbf{x}^i, \mathbf{y}) = \sum_k \phi(\mathbf{x}^i, k, \mathbf{y}_k, \mathbf{y}_{k-1}, \mathbf{y}_{k-2})$$

Keep in mind that depending on how you factorize your feature vector, the details of your viterbi algorithm will change a bit. In this assignment, it is your job to figure out the technical details on your own.

Algorithm 1 Structured Perceptron Algorithm for Sequential Tagging

Input: M tagged sequences $\mathbf{x}^i, \mathbf{y}^i, i = 1 \dots M$ as training data. Number of iterations T .

```

1 Initialization: Randomly initialize parameters  $\mathbf{w}$ , and let  $\bar{\mathbf{w}}$  denote averaged parameters.
2  $count = 0$ 
3 for  $t = 1 \dots T, i = 1 \dots M$  do
4   Predict the current best sequence  $\tilde{\mathbf{y}}$  for  $\mathbf{x}^i$  under current parameters  $\mathbf{w}$ :
                                     
$$\tilde{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} \{ \mathbf{w} \cdot \Phi(\mathbf{x}^i, \mathbf{y}) \} \tag{1}$$

5   if  $\tilde{\mathbf{y}} \neq \mathbf{y}^i$  then
       Update parameters:
                                     
$$\mathbf{w} \leftarrow \mathbf{w} + \Phi(\mathbf{x}^i, \mathbf{y}^i) - \Phi(\mathbf{x}^i, \tilde{\mathbf{y}}) \tag{2}$$

       Update averaged parameters:  $\bar{\mathbf{w}} \leftarrow \bar{\mathbf{w}} + \mathbf{w}$ 
        $count \leftarrow count + 1$ 
6   end
7 end
Output: Averaged parameter vector  $\bar{\mathbf{w}}/count$ 

```

When to stop training? You can determine T based on the performance on the dev set. If you check the performance on the dev set too frequently, it would slow down the learning procedure. Thus, you can check at with some intervals, for example, $T = 1, 5, 10, 15, \dots$

1.2 The CoNLL-2003 NER Dataset

In *conll03_ner.tar.gz*, you will be given the full and the reduced set of the CoNLL-2003 NER corpus [1] [3], feel free to use either of them. The number of sentences and tokens for the dataset are shown in Table 1. There are only four types of named entities in CoNLL-2003: *Person*, *Location*, *Organization* and *Miscellaneous*.

Input Format & Tagging Scheme

The data files are in the following format:

	Full set			Reduced set		
	File	Sentences	Tokens	File	Sentences	Tokens
Training set	<i>eng.train</i>	14,987	203,621	<i>eng.train.small</i>	3,000	41,732
Development set	<i>eng.dev</i>	3,466	51,362	<i>eng.dev.small</i>	500	7,342
Test set	<i>eng.test</i>	3,684	46,435	<i>eng.test.small</i>	500	6,288

Table 1: Number of sentences and tokens in the NER dataset.

```

U.N.      NNP  I-NP  I-ORG
official  NN   I-NP  0
Ekeus     NNP  I-NP  I-PER
heads     VBZ  I-VP  0
for       IN   I-PP  0
Baghdad   NNP  I-NP  I-LOC
.         .    0    0

```

Each line contains four whitespace-separated columns: *word*, *POS tag*, *syntactic chunk tag* and *NER tag*. The first three columns (word, pos-tag, chunk-tag) are part of the input, the fourth column (NER-tag) is what you need to predict. Given that there are four different NER types (*Person*, *Location*, *Organization* and *Miscellaneous*), what is the cardinality of your tag set? Note that both the syntactic chunk tag and the NER tag use BIO encoding. Sentence boundaries are marked by a single blank line. You can find more details about the input/output format here: <http://www.cnts.ua.ac.be/conll2003/ner/>.

Evaluation

For evaluation, you will need to append your predicted tag at the end of each input line, write the sentences into a file (i.e. output.txt), and run the *conlleval.txt* script:

```
./conlleval.txt < output.txt
```

Here is an example of the output file, with the predicted tags in the fifth column:

```

U.N.      NNP  I-NP  I-ORG  0
official  NN   I-NP  0      0
Ekeus     NNP  I-NP  I-PER  I-PER
heads     VBZ  I-VP  0      0
for       IN   I-PP  0      0
Baghdad   NNP  I-NP  I-LOC  I-PER
.         .    0    0      0

```

If your output format is correct, you will see a message like this:

```

processed 7342 tokens with 817 phrases; found: 822 phrases; correct: 644.
accuracy:  96.30%; precision:  78.35%; recall:  78.82%; FB1:  78.58
    LOC: precision:  81.92%; recall:  83.53%; FB1:  82.72  260
    MISC: precision:  82.52%; recall:  76.58%; FB1:  79.44  103
    ORG: precision:  75.45%; recall:  65.28%; FB1:  70.00  167
    PER: precision:  75.34%; recall:  85.27%; FB1:  80.00  292

```

1.3 Deliverables

- (2pt) Define the feature templates you have used for implementing the structured perceptron classifier. Propose 3 additional feature templates that (1) will most likely be useful for improving NER performance but (2) are not part of your implementation.
- (2pt) Write out the viterbi decoding algorithm using one recursive equation.
- (2pts) **Ablation Study:** Ablation study is to remove a subset of features (i.e., a subset of feature templates) and see how that affects the overall performance. Design your own ablation study and report how the performance varies with respect to both the dev set and the test set. Include no more than 4 variations of the feature vector definition, where one of them should be the full model with the complete feature set.
- (2pt) While BIO encoding is the most common, there are a number of other encoding schemes (<https://lingpipe-blog.com/2009/10/14/coding-chunkers-as-taggers-io-bio-bmewo-and-bmewo/>) also used in practice. In fact, some studies have reported that **IO encoding** that does not differentiate 'B' from 'I' can often results in comparable (or even better) performance depending on the data and the task. Why might this be the case? Discuss the potential pros and cons of BIO and IO encoding schemes for NER.
- (2pt) Error analysis and discussion.

1.4 Tips

- The accuracy will depend on the feature vector you design. As a point of reference, the accuracy of competitive systems on the full data should be in the range of 75% - 90% F1 and the accuracy on the reduced data will be in the range of 65%-80% F1. You can find performance of other systems in the CoNLL-2003 summary paper [3] or at the ACL wiki: [http://www.aclweb.org/aclwiki/index.php?title=CONLL-2003_\(State_of_the_art\)](http://www.aclweb.org/aclwiki/index.php?title=CONLL-2003_(State_of_the_art)).
- Using too many features might make training very slow or cause over-fitting. You could limit the number of features by discarding features that occur less than k times in the training set, where typical choice of k could be somewhere between 3 - 5.
- Your feature vectors will be extremely high dimensional with mostly zero entries. Thus arrays shouldn't be your choice of data structure.

References

- [1] Language-independent named entity recognition (ii). <http://www.cnts.ua.ac.be/conll2003/ner/>.
- [2] Michael Collins. Discriminative training methods for hidden markov models: Theory and experiments with perceptron algorithms. In *Proceedings of the ACL-02 conference on Empirical methods in natural language processing-Volume 10*, pages 1–8. Association for Computational Linguistics, 2002.
- [3] Erik F Tjong Kim Sang and Fien De Meulder. Introduction to the conll-2003 shared task: Language-independent named entity recognition. In *Proceedings of the seventh conference on Natural language learning at HLT-NAACL 2003-Volume 4*, pages 142–147. Association for Computational Linguistics, 2003.
- [4] Maksim Tkachenko and Andrey Simanovsky. Named entity recognition: Exploring features. 2012.